

Data access layer optimisation of the Gaia data processing in Barcelona for spatially arranged data

Master in Innovation and Research in Informatics

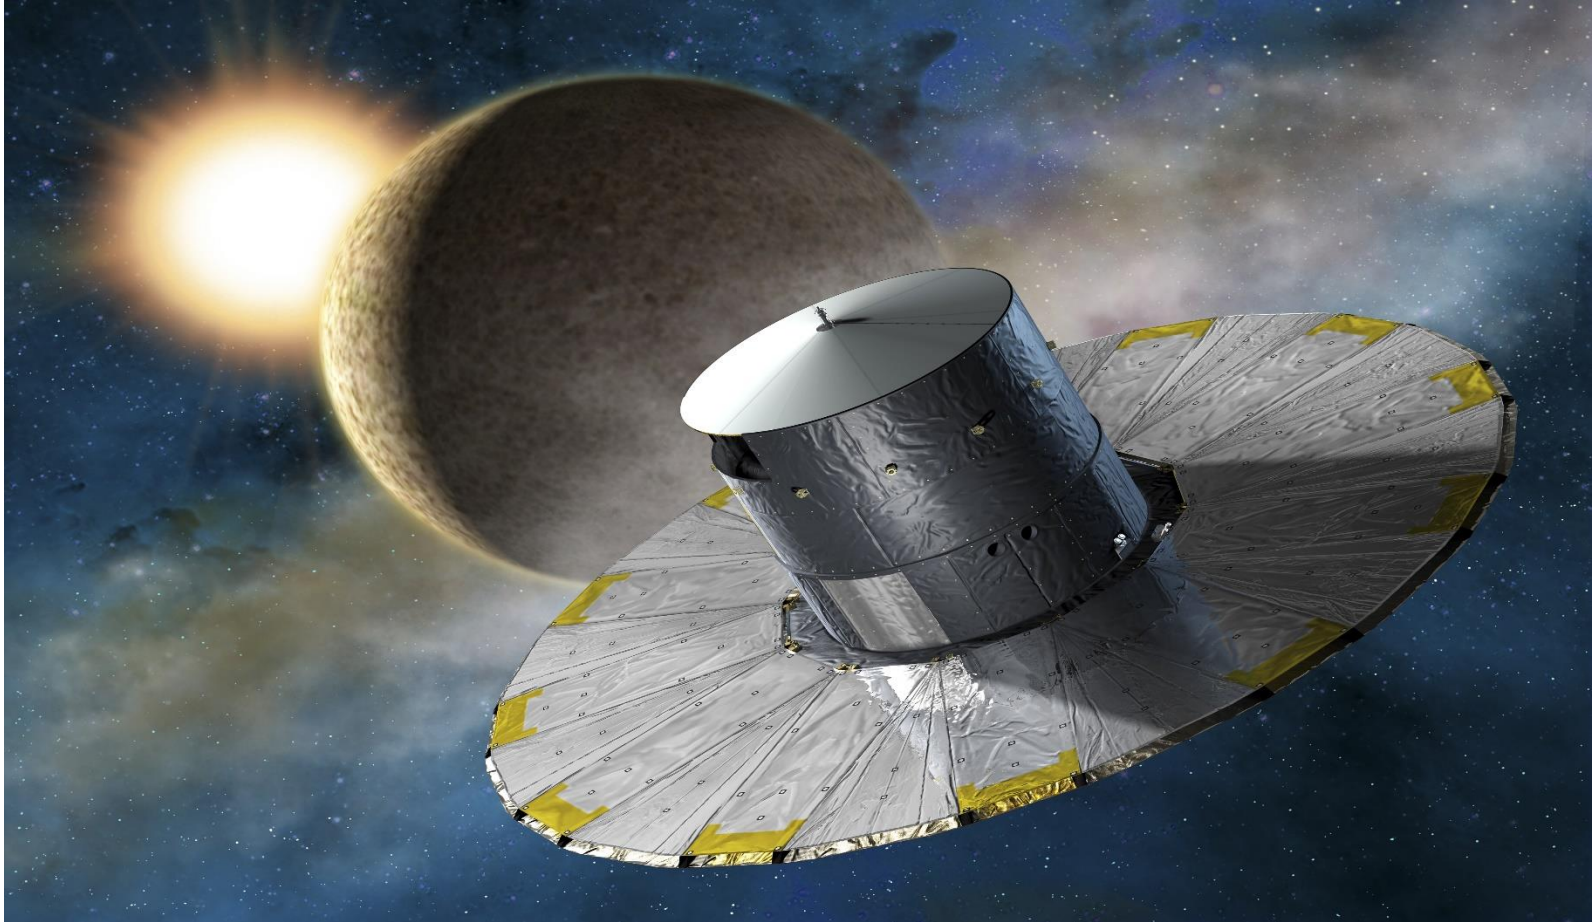
Francesca Ronchini

E-mail: ronchinifrancesca8@gmail.com

Outline

1. The Gaia mission
2. HEALPix sphere tessellation
3. HDF5 file format
4. Implementation
5. Test and results
6. Conclusions

The Gaia mission

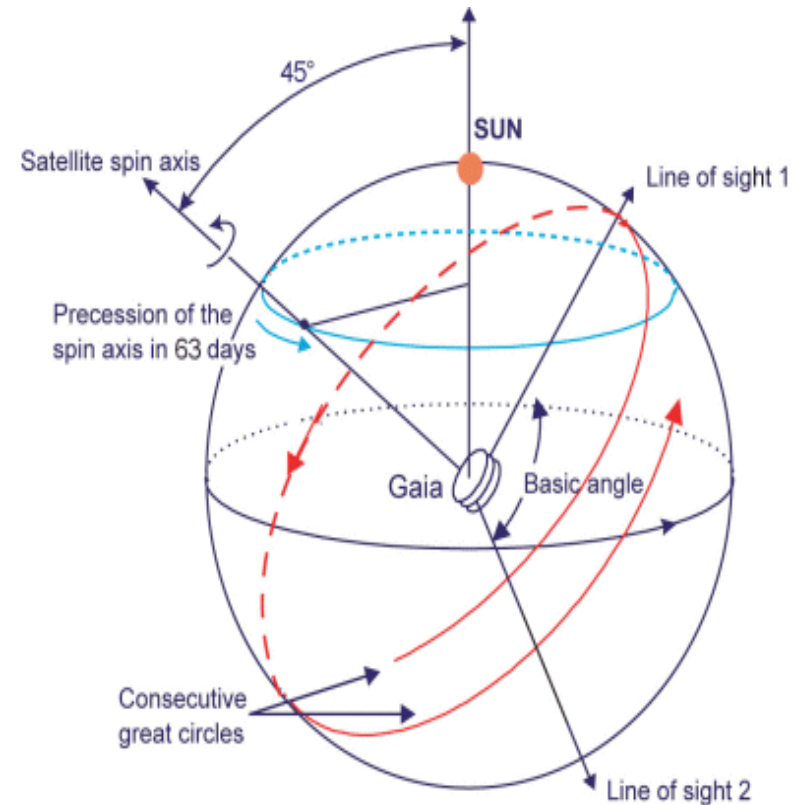


The Gaia mission

One of the most challenging
European Space Agency (ESA)
space missions.

It was launched on **December
19th, 2013**, from Kourou, in
French Guiana.

It operates at the **L2 point of the
Solar System**, about 1.5 million
kilometers from the Earth.



The Gaia mission

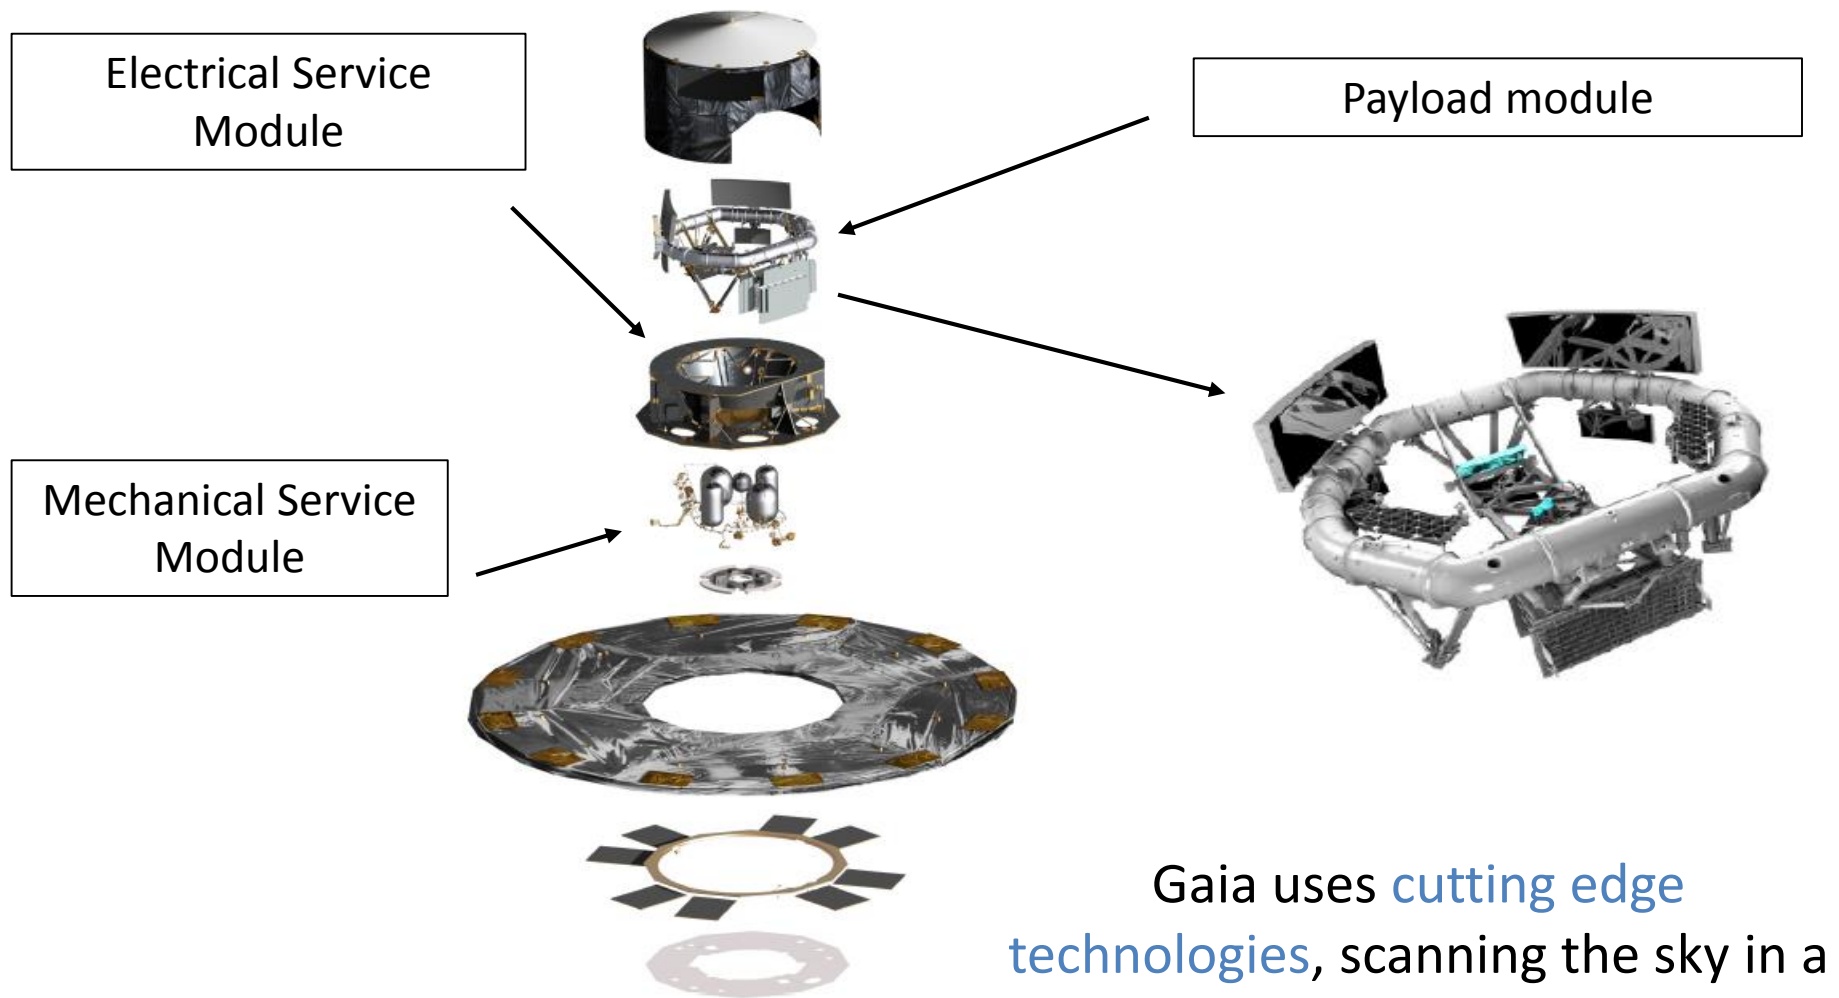
The satellite will observe each source about 85 times on average, recording:

- Colour
- Brightness
- Position

Scope of the mission: **the most precise three-dimensional map of our Galaxy**, the Milky Way.

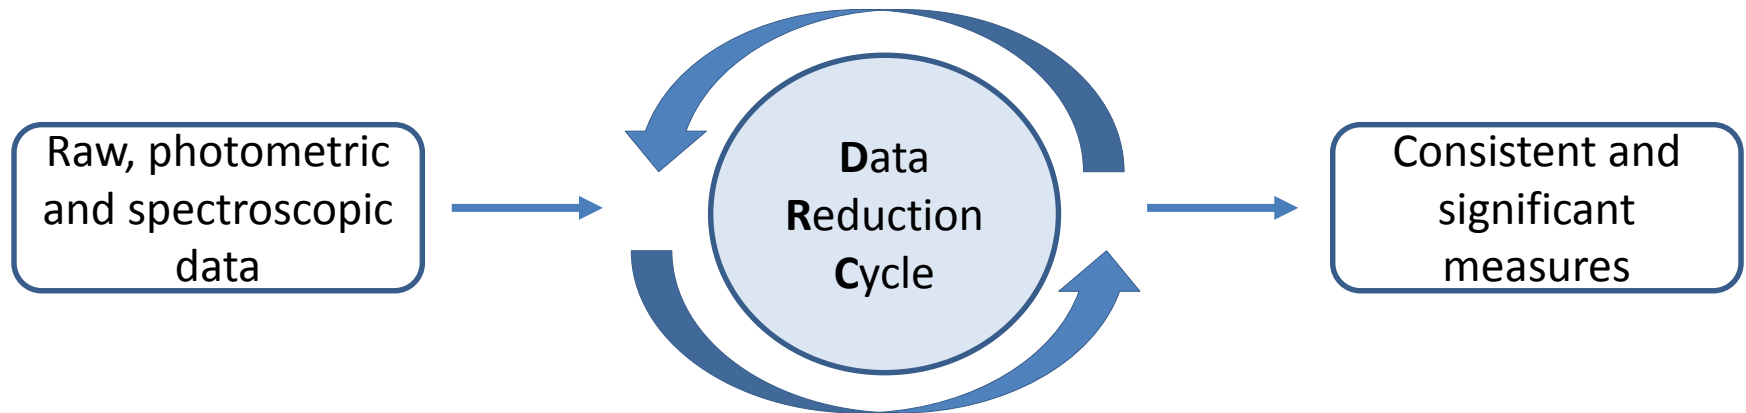
Duration of the mission: **5 years**

The spacecraft



Gaia uses **cutting edge technologies**, scanning the sky in a continuous way.

Gaia data processing



DPAC Consortium is in charge to process data.

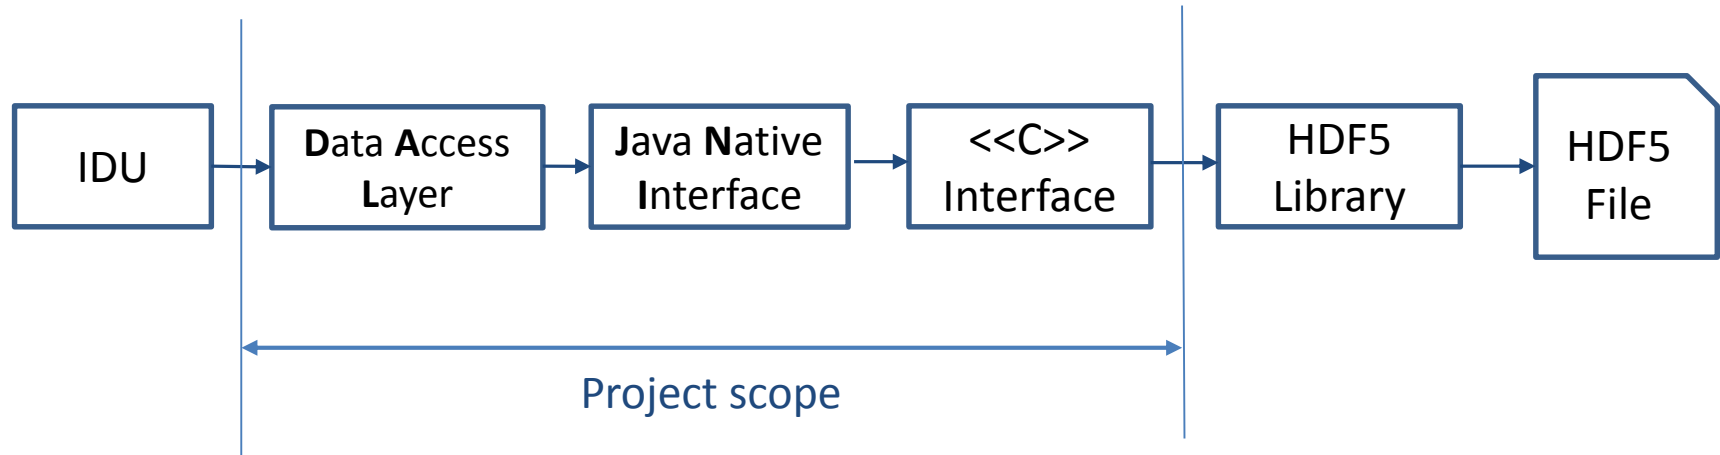
Intermediate **Data Updating**:

- **Cross-Match**: provides links between the individual detections and the source entries in the catalogue.

Challenges:

- Number of detections to handle: **10^{11} ! Impossible in a single process!**
- Data organization

The scope of the project



Constraints:

- **No database** can be used because of the Marenostrum III supercomputer. All I/O operations must be based on files.
- Gaia data is stored as **Zip-compressed serialized Java objects**.
- The frontend software and interfaces must be in **Java**

HEALPix sphere tessellation

The library allows:

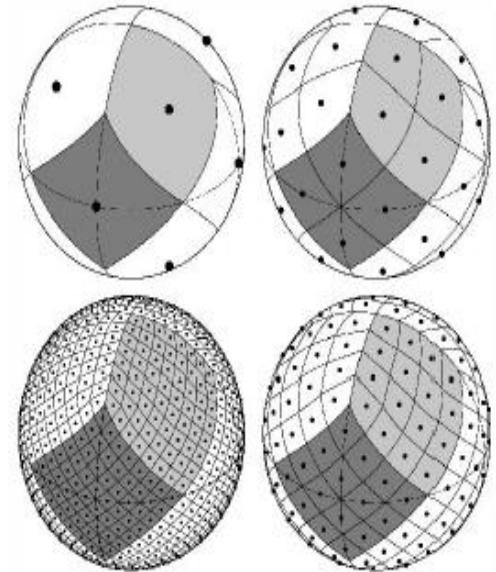
- To partition and reduce the data coming from spherical topologies in a fast and efficient way, especially in cases as IDU
- To refer to the data distributed on the sphere in a hierarchical way
- To access the data in a more direct way

The library was developed for the cosmic microwave background experiments in the 1990s.

HEALPix sphere tessellation

Main Features:

- Each pixel contains the **same** surface of area
- 12 pixels in 3 rings around the poles and equator
- For each level: **4 new pixels**, coded adding one digit between 0 and 3 to the previous base pixel



Tessellation of the spherical surface from level 0 to 3

HDF5 file format

Powerful set of open-source libraries that allows to save, organize and manage data when it is not possible, or it is restrictive, to use traditional database systems.

Main features:

- It is **portable**
- It allows to refer to **hierarchical data objects** in a natural manner
- It allows both **sequential and coordinate-style lookup**

It was developed for the first time by the National Center for Supercomputing Applications (NCSA) in 1987. Today it is supported by the HDF Group, University of Illinois.

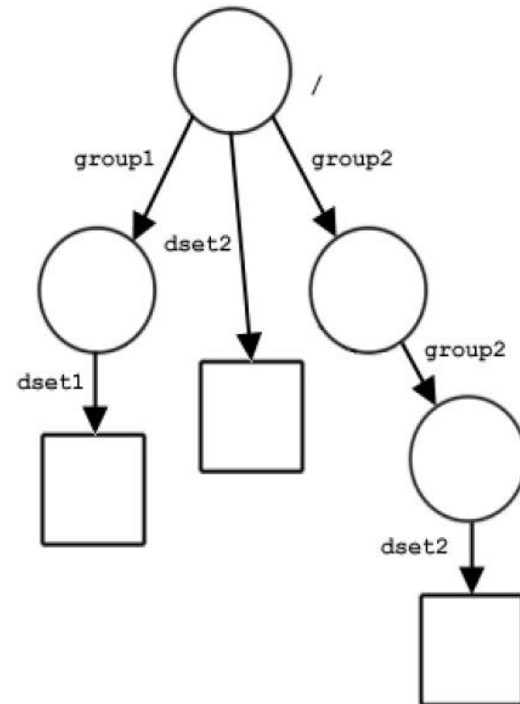
HDF5 file format

Container of two primary type of objects:

- **Groups**: they contain other groups or objects
- **Datasets**: they contain data written in a specific format

Other types of objects:

- Dataspaces
- Datatype
- Attributes



HDF5 file format

Gaia Binary file format (GBIN)

- It is not adapted for data processing
- It only allows to access the data sequentially without providing any structure within the file.
- It can not provide any of the required data organization within the file. It is required offline data arrangement
- Difficult to scale and distribute

HDF5

- It allows the implementation of hierarchical structure file
- Clear structure of the files
- It is not mandatory to specify the type of the data, using the **opaque** datatype
- Direct access to the data
- File self-described
- Binary data can be included
- No limit on the quantity of data

Implementation

Two solutions:

- Full hierarchical HEALPix structure
- Flat HEALPix structure

Dataset numbering:

- Numbering using **counters**
- Numbering using **attributes**

Java Native Interface

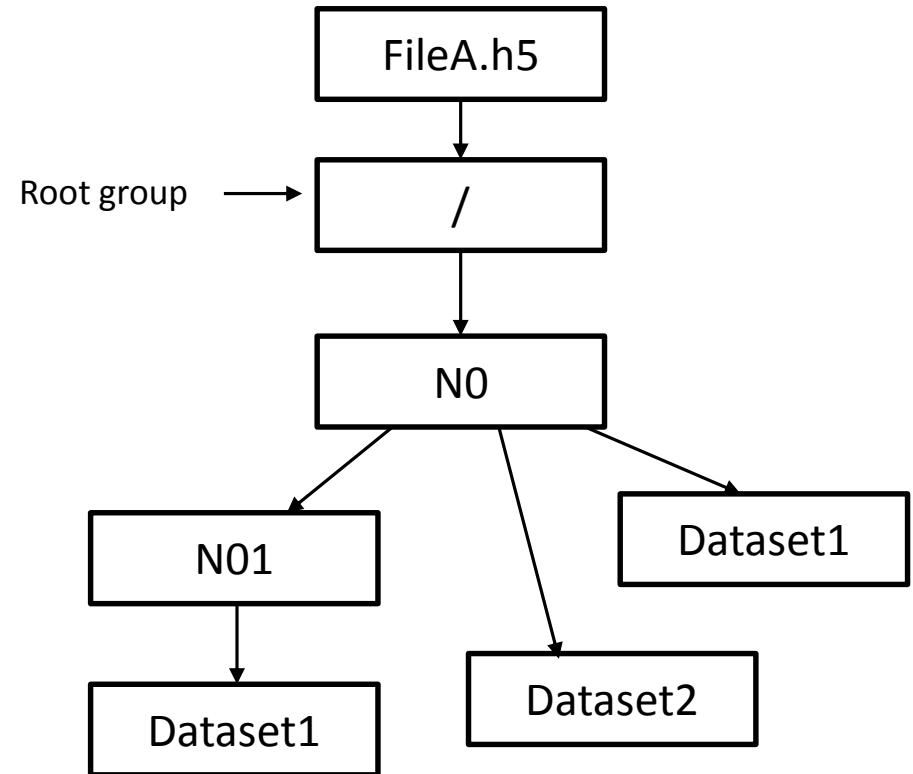
DpcbTools interface

Full hierarchical HEALPix structure

Hierarchical **tree structure** following the HEALPix structure

Groups accessible through full or relative paths:

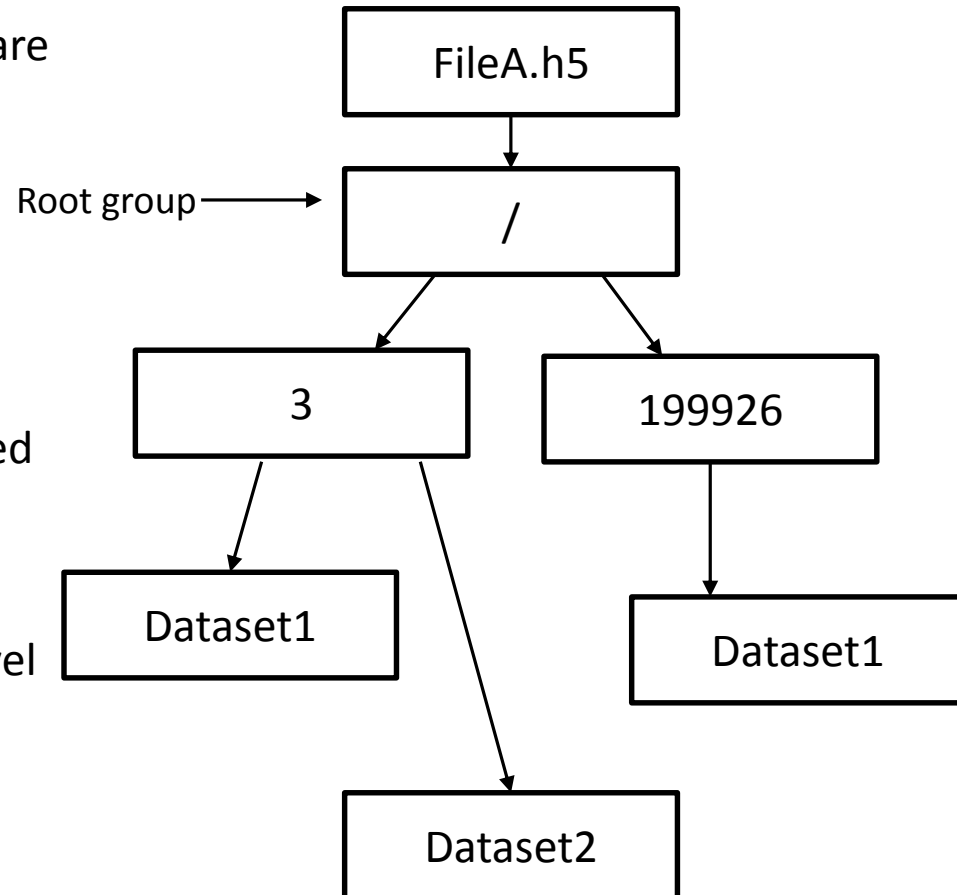
- **/N0/N01/Dataset1**: full path to Dataset1
- **N01/Dataset1**: relative path, from N0, to Dataset1



Flat HEALPix structure

All groups are at the HDF5 root level and are identified by a **unique** HEALPix numerical identifier

- Pixel range: $[0; 12(N_{\text{size}})^2 - 1]$, where:
 - $N_{\text{size}} = 2^k$.
 - k : level at which the data are considered
- For this project:
 - $[0; 201326591]$, as 12 is the highest level



Dataset numbering

Two alternatives

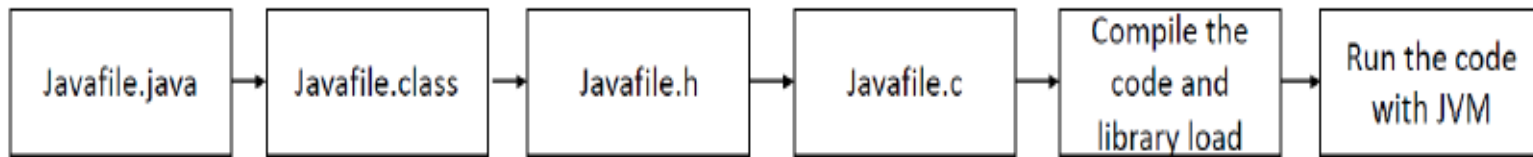
Counters

- Linear and straightforward
- First dataset: Dataset1
- The group already exists: using the H5L interface and updating a counter variable

Attributes

- To store meta-data
- First dataset: Dataset1 and NUM_OF_DATASET_ATTR
- The group already exists: using the H5A interface and updating the attribute

Java Native Interface



- Prototype of the function: *Java_package_and_classname_function_name*(JNI arguments)
- They have to be loaded and may return values back to the Java VM

DpcbTools interface

The Java part of this project has been integrated into a software library named **DpcbTools**, in coordination with the [Barcelona IDU/DPCB team](#).

FileHdf5Archiver

FileHdf5Archiver
<ul style="list-style-type: none">- fileId: Integer- datasetId: Integer- nat: Hdf5ArchiveNative- archiveStructure: Hdf5ArchiveStructure
<ul style="list-style-type: none">+ FileHdf5Archiver()+ FileHdf5Archiver(h5File: File)+ add(bos: ByteArrayOutputStream, entry: String): void+ add(file: File, entry: String): void+ add(data: List<? extends GaiaRoot>, entry: String, format: FileFormat): void+ setHealpixStructure(structure: Hdf5ArchiveStructure): void+ writeHealpixFlat(HealpixName: int, data: byte[], Level: int): int+ writeHealpixHierarchical(HealpixName: String, data: byte[], Level: int): int

FileArchiveHdf5FileReader

FileArchiveHdf5FileReader
<ul style="list-style-type: none">- fileId: int- healpixName: String- healpixIndex: int- readObjects: long- totalObjects: long- datasets: List<String>- nat: Hdf5ArchiveNative- flatStructure: boolean
<ul style="list-style-type: none">+ FileArchiveHdf5FileReader(archive: File, entryName: String)+ getAllObjects(): List<E>+ getCount(): long+ getNameList(healpixName: String): List<String>+ getNameListFlat(): List<String>+ getNextObject(): E+ getRemainingObjects(): List<E>+ hasNext(): boolean

Results

Test the feasibility and performance when creating the complete HDF5 file structures

HDF5 Hierarchical file format

Level	Creation time [s]	File Size	Number of groups	Number of Healpix pixels
0	0.012	10 KB	12	12
1	0.017	48 KB	60	48
2	0.013	199 KB	252	192
3	0.028	802 KB	1020	768
4	0.093	3 MB	4092	3072
5	0.378	13 MB	16380	12288
6	1.654	52 MB	65532	49152
7	16.563	206 MB	262140	196608
8	27.610	824 MB	1048572	786432
9	233.954	3 GB	4194300	3145728
10	461.757	13 GB	16777212	12582912
11	N/A	N/A	67108860	50331648
12	N/A	N/A	268435452	201326592

HDF5 Flat file format

Level	Creation time [s]	File Size	Number of groups	Number of Healpix pixels
0	0.018	10 KB	12	12
1	0.014	39 KB	48	48
2	0.020	155 KB	492	492
3	0.027	615 KB	768	768
4	0.075	2 MB	3072	3072
5	0.253	10 MB	12288	12288
6	1.177	39 MB	49152	49152
7	4.731	157 MB	196608	196608
8	22.163	629 MB	786432	786432
9	12125.837	3GB	3145728	3145728
10	N/A	N/A	12582912	12582912
11	N/A	N/A	50331648	50331648
12	N/A	N/A	201326592	201326592

Results

- File format considered: HDF5-based alternatives, TAR and ZIP
- Set of data considered: AstroObservation and Match
- File system considered: GPFS and local 500 GB IBM hard drive
- Iterations executed: 3 times for each test
- Size of set of data: 1GB

Results

Level	TAR	ZIP	HDF5_H	HDF5_F
0	22.05	16.98	21.77	22.50
2	21.84	17.04	21.92	22.43
4	22.28	17.06	21.98	22.41
6	22.51	17.21	22.43	16.37

Write results on the GPFS file system for AstroObservation data in [MiB/s]

Level	TAR	ZIP	HDF5_H	HDF5_F
0	22.05	16.98	21.77	22.50
2	21.84	17.04	21.92	22.43
4	22.28	17.06	21.98	22.41
6	22.51	17.21	22.43	16.37

Write results on the GPFS file system for Match data in [MiB/s]

Results - Global

Write Level	Read Level	TAR	ZIP	HDF5_H	HDF5_F
0	0	81.52	83.55	147.96	146.66
	2	24.57	26.44	46.14	45.56
	4	5.38	5.65	10.05	10.05
	6	0.96	1.00	1.74	1.74
2	0	68.42	82.87	143.07	141.16
	2	53.63	81.57	128.31	137.07
	4	12.38	17.96	30.69	30.51
	6	2.07	2.79	4.83	4.82
4	0	47.71	81.81	132.27	96.09
	2	39.33	80.74	121.83	95.84
	4	21.73	76.25	77.59	70.77
	6	3.92	10.68	12.58	11.38
6	0	15.75	69.79	25.12	8.48
	2	14.45	67.12	18.27	7.88
	4	11.42	62.15	8.00	5.15
	6	4.94	49.17	1.94	1.81

Read results on the GPFS file system
in [MiB/s] for AstroObservation data

Write Level	Read Level	TAR	ZIP	HDF5_H	HDF5_F
0	0	47.78	50.40	51.66	54.25
	2	48.05	55.26	53.40	63.86
	4	7.75	8.03	8.73	8.84
	6	0.51	0.53	0.57	0.57
2	0	47.95	50.91	54.14	56.15
	2	48.40	50.99	57.74	53.05
	4	7.81	7.99	8.48	8.79
	6	0.52	0.53	0.58	0.78
4	0	47.12	44.78	61.55	63.39
	2	45.81	54.76	62.62	65.43
	4	34.92	44.05	60.18	61.36
	6	2.84	3.79	4.63	4.65
6	0	20.22	53.06	59.23	38.99
	2	20.58	49.50	58.93	38.52
	4	18.97	48.64	54.55	34.17
	6	8.26	45.49	27.09	21.07

Read results on the GPFS file system
in [MiB/s] for Match data

Results - Single

Write Level	Read Level	TAR	ZIP	HDF5_H	HDF5_F
0	0	68.94	81.23	136.70	142.72
	2	17.63	20.66	36.50	36.34
	4	3.82	4.40	7.47	7.76
	6	0.71	0.83	1.44	1.44
2	0	55.62	80.53	136.90	133.95
	2	37.35	71.50	108.01	108.08
	4	7.98	15.19	24.19	23.84
	6	1.50	2.88	4.47	4.39
4	0	37.81	76.37	117.88	78.05
	2	27.25	68.52	82.63	60.93
	4	14.82	53.84	74.09	42.43
	6	2.84	10.09	8.84	7.91
6	0	15.80	58.45	17.35	7.45
	2	12.76	51.26	12.09	5.67
	4	8.42	41.39	6.22	3.85
	6	4.28	29.80	1.95	1.82

Read results on the GPFS file system
in [MiB/s] for AstroObservation data

Write Level	Read Level	TAR	ZIP	HDF5_H	HDF5_F
0	0	53.50	59.49	75.30	76.33
	2	54.80	58.38	75.38	77.64
	4	11.87	13.36	16.57	16.27
	6	0.80	0.86	1.04	1.43
2	0	53.71	58.61	75.24	77.12
	2	53.69	59.42	75.84	75.93
	4	11.72	13.33	16.12	16.09
	6	0.79	0.86	1.04	1.04
4	0	51.07	58.13	77.17	76.76
	2	51.57	58.03	72.76	74.34
	4	45.33	59.09	47.33	75.88
	6	2.90	3.91	1.09	5.03
6	0	25.82	57.84	64.77	47.09
	2	26.06	58.08	61.01	45.50
	4	21.08	57.81	57.92	39.69
	6	9.57	54.50	29.92	25.66

Read results on the GPFS file system
in [MiB/s] for Match data

Conclusions

- In many cases the performance is increased by more than 50%.
- The work presented has substantially contributed to the optimization of the DAL
- The work will help to collect and store the data in a very more efficient way.
- The data can now be stored according to their spatial distribution in fewer well structured

Future work:

- To explore the native HDF5 compression algorithms
- To include time information within the spatially distributed data structure

Thank you